Decision trees and linear regression with PYTHON

Objective: The objective of this exercise is to become familiar with fitting decision trees and linear regression models in Python.

Material: Lecture notes *"Introduction to Machine Learning and Data Mining"* as well as the files provided from DTU Learn.

5.1 Decision Trees

In this part of the exercise we will fit decision trees using the Python class `sklearn.tree.DecisionTreeClassifier`.

As a splitting criterion, the function uses one of the two impurity measures:

$$\text{gdi}(t) = 1 - \sum_{i=0}^{c-1} p(i|t)^2 \qquad \text{equivalent to Gini}(t),$$

$$\text{deviance}(t) = -2 \sum_{i=0}^{c-1} p(i|t) \log p(i|t) \quad \text{equivalent to Entropy}(t).$$

We will analyze the vertebrate data given in Table 4.1 of "Introduction to Data Mining" [1] as well as the wine data we have used in the previous exercise (this data is similar to the data introduced in the chapter on "Tree-based methods" in the book "Introduction to Machine Learning and Data Mining").

The vertebrate data set has the following attributes, which are all categorical:

Vertebrate dataset

#	Attribute	Unit
1	Body temperature	0:Cold blooded, 1:Warm blooded.
2	Skin cover	0:None, 1:Hair, 2:Scales, 3:Feathers, 4:Fur, 5:Quills
3	Gives birth	0:No, 1:Yes.
4	Aquatic creature	0:No, 1:Yes, 2:Semi.
5	Aerial creature	0:No, 1:Yes.
6	Has legs	0:No, 1:Yes.
7	Hibernates	0:No, 1:Yes.

The wine data set have the following attributes, where 1–11 are continuous and 12 is categorical:

Wine dataset

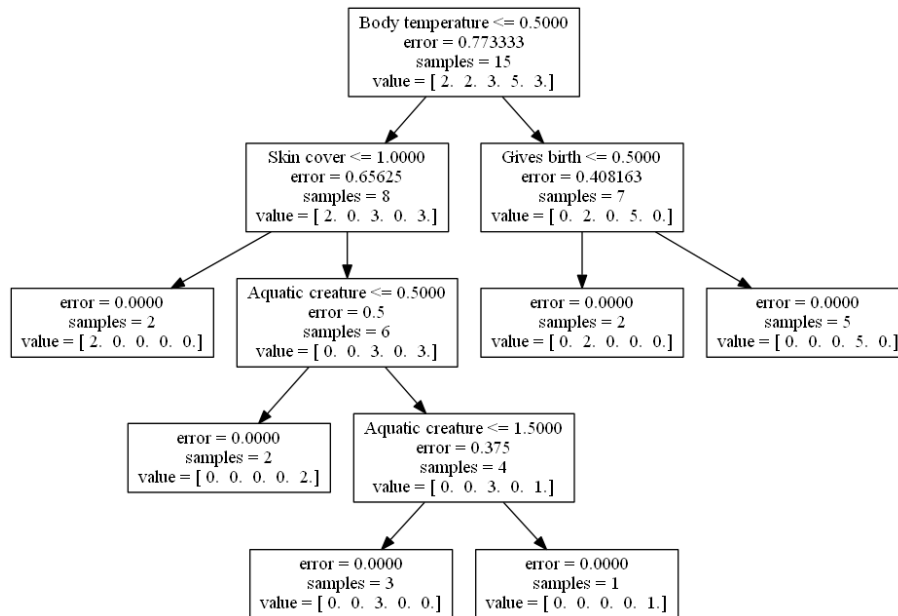| #	Attribute	Unit
1	Fixed acidity (tartaric)	g/dm^3
2	Volatile acidity (acetic)	g/dm^3
3	Citric acid	g/dm^3
4	Residual sugar	g/dm^3
5	Chlorides	g/dm^3
6	Free sulfur dioxide	mg/dm^3
7	Total sulfur dioxide	mg/dm^3
8	Density	g/cm^3
9	pH	pH
10	Sulphates	g/dm^3
11	Alcohol	% vol.
12	Quality score	0–10

Figure 1: Decision tree of the vertebrates data

5.1.1 Examine and run the script `ex5_1_1.py`, which contains the vertebrates data given in Table 4.1 of "Introduction to Data Mining".

5.1.2 Examine and run the script `ex5_1_2.py`. The script fits a decision tree to the vertebrates data using the Gini splitting criterion and stops splitting only when nodes are pure.

Script details:

- *Import module `tree` from `sklearn` and type `help(tree.DecisionTreeClassifier)` to learn how to fit a decision tree in Python.*

- *The parameter `criterion` can be used to choose the splitting criterion (Gini or Entropy).*

- *The parameters `min_samples_split`, `min_samples_leaf`, and `max_depth` influence the stopping criterion.*

- *The parameter `feature_names` can be used during the exporting to supply the names of the attributes.*

Did you get a similar result to the figure 1? Try to understand the information provided by the tree.

5.1.3 Examine and run the script `ex5_1_3.py`. The script shows how the tree changes when you use the Entropy (i.e. `deviance`) splitting criterion instead of the Gini (`gdi`).

Script details:

- *The parameter `criterion` can be used to choose the splitting criterion used.*

5.1.4 Use the script `ex5_1_4.py` to show that a dragon, which is cold-blooded, has scales, gives birth, is semi aquatic, is an aerial creature, has legs, and hibernates, will be classified as a reptile.

Script details:

> · *After fitting decision tree, you can use the method `predict()` of the tree object to classify a new data object.*

5.1.5 Load the wine data set into Python using `ex5_1_5.py`.

The script removes outliers from the data set. Outliers are defined as data objects where Volatile acidity > 20 g/dm^3, Density > 10 g/cm^3, or Alcohol $> 200\%$.

The script also removes attribute 12, Quality score, from the data set, in order to use only the quantitative measurements for the prediction.

Script details:

> · *You have worked with the wine data before in exercise 4.2.1 You can use your solution to that exercise as a starting point.*
>
> · *To select a subset of data objects from the data matrix, you can write `X=X[n,:]` where `n` is either a vector of indices of the data objects to be selected, or a binary mask. Remember to update the class index vector accordingly `y=y[n,:]`, and to recompute `N`.*
>
> · *Similarly you can select a subset of the attributes from the data matrix `X=X[:,m]` where `m` is either a vector of indices to the attributes, or a binary mask. Remember to update the list of attribute names accordingly, and to recompute the number of attributes `M`.*

You should end up with a data matrix of size $N \times M = 6304 \times 11$.

5.1.6 Inspect and run the script `ex5_1_6.py`. The script fits a decision tree to the wine data in order to estimate if the wine is red or white. The script uses the Gini splitting criterion and the following parameter value as stopping criterion `min_samples_split=100`.

Script details:

> · *The solution to exercise 5.1.2 is used as a starting point.*

Explain what happens when you change the values of the parameter `min_samples_split`.

Observe the changes in the tree size after reducing the parameter.

5.1.7 Explain how the script `ex5_1_7.py` shows that a wine with the following attribute values would be classified as White.

#	Attribute	Value
1	Fixed acidity (tartaric)	6.9
2	Volatile acidity (acetic)	1.09
3	Citric acid	0.06
4	Residual sugar	2.1
5	Chlorides	0.0061
6	Free sulfur dioxide	12
7	Total sulfur dioxide	31
8	Density	0.99
9	pH	3.5
10	Sulphates	0.44
11	Alcohol	12

5.2 Linear and logistic regression

Regression is widely used in datamining and machine learning for predicting an output (dependent variable) for a data object given its attributes. Regression can be defined as the task of learning a target function that maps a set of attributes onto an output.

We will initially consider a linear regression problem where we will model the outputs, $\boldsymbol{y} = [y_1, y_2, \ldots, y_N]^\top$, based on a single attribute for each data object, $\boldsymbol{X} = [x_1, x_2, \ldots, x_N]^\top$, using the model

$$y_n = w_0 + w_1 x_n + \epsilon_n,$$

where ϵ_n is residual noise, w_0 a constant offset and w_1 the linear coefficient for x_n. We can write an equation for each data object,

$$
\begin{aligned}
y_1 &= w_0 + w_1 x_1 + \epsilon_1 \\
y_2 &= w_0 + w_1 x_2 + \epsilon_2 \\
&\vdots \\
y_N &= w_0 + w_1 x_N + \epsilon_N.
\end{aligned}
$$

Introducing a vector of all ones $\mathbf{1} = [1, 1, \ldots, 1]^\top$ we can write the N equation compactly in matrix-vector form

$$\boldsymbol{y} = w_0 \mathbf{1} + w_1 \boldsymbol{x} + \boldsymbol{\epsilon} = \begin{bmatrix} \mathbf{1} & \boldsymbol{x} \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \end{bmatrix} + \boldsymbol{\epsilon}.$$

In order to estimate w_0 and w_1 we minimize the least squares error function, i.e.

$$(w_0, w_1) = \arg\min_{w_0, w_1} \sum_{n=1}^{N} (y_n - (w_0 + w_1 x_n))^2 = \arg\min_{w_0, w_1} \left\| \boldsymbol{y} - \begin{bmatrix} \mathbf{1} & \boldsymbol{x} \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \end{bmatrix} \right\|_F^2.$$

Based on our linear regression model we can predict the value of an unobserved output y^* from an observed measurement x^* as

$$y^* = w_0 + w_1 x^*.$$

5.2.1 Inspect and run the script `ex5_2_1.py`. The script generates 100 data objects according to the linear regression model,

$$y_n = w_0 + w_1 x_n + \epsilon_n,$$

Each of the n observations (having one attribute value x_n) should be set to $x_1{=}0, x_2{=}1, x_3{=}2, \ldots, x_N{=}N{-}1$, and ϵ_n is generated at random from a Normal distribution with mean zero and standard deviation 0.1. Inspect the script and verify the parameters of the model are $w_0 = -0.5$ and $w_1 = 0.01$.

Script details:

· *You can use functions* `range()` *or* `linspace()` *to make a sequence of integer or real numbers.*

· *You can generate Normal distributed random numbers using the function* `random.randn()`.

Make a scatter plot of your data. Does it look like the plot in figure 2?

5.2.2 Inspect and run the script `ex5_2_2.py`. Explain how the script estimates the model parameters, w_0 and w_1, from the data you have generated. Show that the estimated parameters are close but not exactly equal to the parameters used to generate the data.

Notice how the script plots predictions of the model as well as the predictions you get using the parameters used to generate the data in the same graph as the scatter plot you made in exercise 5.2.1.
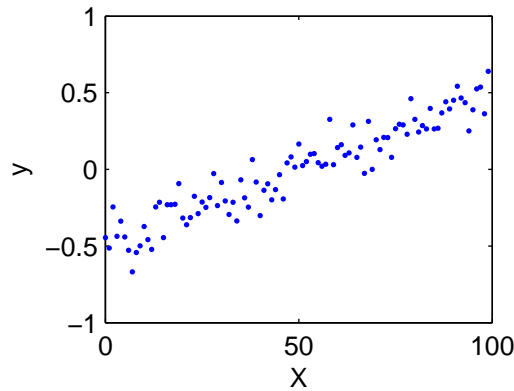
Script details:

Figure 2: Scatter plot of data

- · *You can use the class* `sklearn.linear_model.LinearRegression` *as your model. Set* `fit_intercept` *parameter to True, unless you have added column of "1"s to your data matrix X.*
- · *Use method* `fit()` *to estimate parameters in the linear model, and method* `predict()` *to predict values for new data points.*

Show that the estimated model parameters get closer to the parameters used to generate the data if you include more data objects or reduce the noise variance.

Sometimes there are more complex relationships between the attributes and output. In order to capture more complex relationships we can include attributes that are specified transformations of existing attributes such as $\sqrt{x}$, $\log(x)$, x^2 and x^3. We will presently consider the model

$$y_n = w_0 + w_1 x_n + w_2 x_n^2 + \epsilon_n$$

The above model can again be written for all N observations simultaneously in terms of vectors and matrices, i.e.

$$\boldsymbol{y} = [\mathbf{1} \ \boldsymbol{x}_1 \ \boldsymbol{x}_2] \begin{bmatrix} w_0 \\ w_1 \\ w_2 \end{bmatrix} + \boldsymbol{\epsilon} = [\mathbf{1} \ \boldsymbol{X}]\boldsymbol{w} + \boldsymbol{\epsilon}$$

where $\boldsymbol{x}_1 = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_N \end{bmatrix}$, $\boldsymbol{x}_2 = \begin{bmatrix} x_1^2 \\ x_2^2 \\ \vdots \\ x_N^2 \end{bmatrix}$ and $\boldsymbol{w} = \begin{bmatrix} w_0 \\ w_1 \\ w_2 \end{bmatrix}$.

5.2.3 Inspect and run the script `ex5_2_3.py` which generates a data set according to the linear regression model including the squared term, fits the model

and makes a plot as in the previous exercise. The script allows you to specify separately how many terms to include (squared, cubed, etc.) for generating the data and fitting the model. Illustrate what happens when fitting a model with more higher order terms? How well is the model able to fit the data?

We will return to the wine data and attempt to predict the alcohol content using linear regression. The linear regression model we consider is

$$y_n = w_0 + \sum_k w_k x_{n,k} + \epsilon_n,$$

where y_n is the observed alcohol content, ϵ_n the model residual, w_k the value of the estimated model coefficient for the k^{th} attribute whereas $x_{n,k}$ denotes the n^{th} observation of the k^{th} attribute. As such, $x_{2,1}$ denotes the second observation of the first attribute (Fixed acidity) and $x_{5,2}$ the fifth observation of the second attribute (Volatile acidity).

5.2.4 Inspect and run the script `ex5_2_4.py`. The script loads the Wine data into Python using the solution to exercise 5.1.5. The script also fits a linear model that predicts the Alcohol attribute based on the 10 remaining attributes (again excluding Quality score) and makes a scatter plot of the true versus the predicted alcohol content. Explain how the visualizations show that the alcohol content can be predicted with an accuracy of approximately ±1 percentage point.

Script details:

· *You can take similar approach as in exercises 5.2.2 and 5.2.3 to fit linear regression model and use it to predict values.*

Show from the estimated parameters of the linear regression model that wines with higher alcohol content in general have lower density.

5.2.5 Inspect and run the script `ex5_2_5.py`. The script introduces more variables in the regression by combining or transforming existing variables. Explain how that is done in the script and which transformations are used. Try to introduce your own transformations.

Script details:

· *For example you can include the square of the Fixed acidity in the model and the square of the Volatile as well as a multiplicative interaction between Fixed acidity and Volatile acidity. You can use* `numpy.power()` *and* `numpy.multiply()` *functions on appropriate columns of data matrix.*

· *Function* `numpy.bmat()` *is useful to concatenate matrices and/or vectors, for instance:* `X=np.bmat('X, Xfa`

· *Plots of attributes versus the model residuals can reveal if there is structure in the output that can be explained by transformations of the attributes.*

How low can you get the mean squared error?

Until now we have discussed the most simple linear model based on the squared error cost function. In the generalized linear model, we can also use other cost functions as well as non-linear link functions. For instance, for outputs that are purely binary, $y \in \{0, 1\}$, the logistic function is very suitable, logistic$(s) = \frac{1}{1+exp(-s)}$, which maps the output in the regression analysis to the range (0,1). In figure 3 there is a plot of the logistic link function.

It can be very useful to limit the possible outputs to the range (0,1) if the outputs we wish to predict are binary. Predicting binary outputs is what we have previously referred to as classification, and generalized linear models can thus both be used for regression and classification. Classification with a binomial cost function and logistic link function is called *logistic regression*, and in Python is implemented in module:
`sklearn.linear_model.logistic.LogisticRegression`

We will return to the wine data and predict the type of wine using logistic regression. The model we consider is

$$y_n = \text{logistic} \left(w_0 + \sum_k w_k x_{n,k} \right) + \epsilon_n$$

where y_n is the observed type of wine (Red: $y_n = 0$ and White: $y_n = 1$). We will say that a wine is White with probability p and Red with probability $1 - p$. where $p = \text{logistic}(w_0 + \sum_k w_k x_{n,k})$.

5.2.6 Inspect and run the script `ex5_2_6.py`. The script loads the Wine data into Python using the solution to exercise 5.1.5 and fits a logistic regression model that predicts the type of wine (Red or White) based on the attributes 1-11 (again excluding Quality score). Explain how this is done.
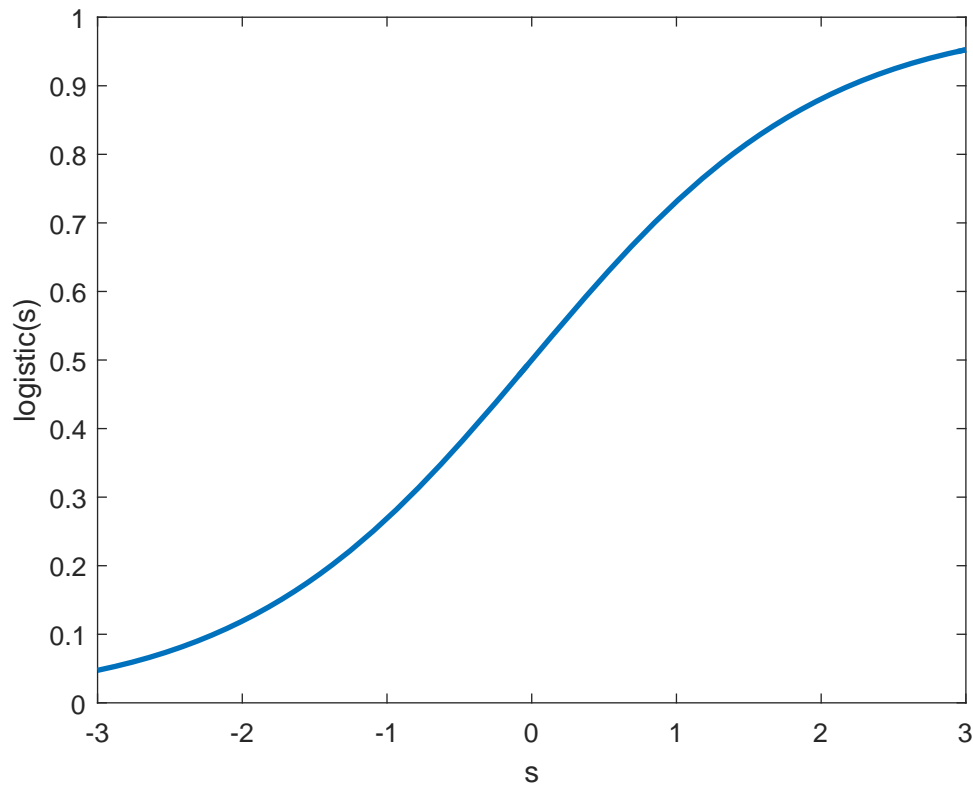
Figure 3: Plot of the logistic link function.

Show that the wine defined in exercise 5.1.7 would be classified as White with more than 90% probability.

Script details:

· *To estimate a logistic regression model, use the class:* `sklearn.linear_model.logistic.LogisticRegression`.

· *To make a prediction from the estimated model, use the function* `predict()`. *If you want to get probabilities of new data object belonging to given class use* `predict_proba()`.

5.3 Tasks for the report

The report will contain three sections, two about regression and one about classification. For the regression section, you will now be able to address this question:

Regression, part a: In this section, you are to solve a relevant regression problem for your data and statistically evaluate the result. We will begin by examining the most elementary model, namely linear regression.

- Explain what variable is predicted based on which other variables and what you hope to accomplish by the regression. Mention your feature transformation choices such as one-of-K coding. Since we will use regularization momentarily, apply a feature transformation to your data matrix $\mathbf{X}$ such that each column has mean 0 and standard deviation 1[1].

Meanwhile, for the **classification** section:

- Explain which classification problem you have chosen to solve. Is it a multi-class or binary classification problem?

---

[1]We treat feature transformations and linear regression in a very condensed manner in this course. Note for real-life applications, it may be a good idea to consider interaction terms and the last category in a one-of-K coding is redundant (you can perhaps convince yourself why). We consider this out of the scope for this report

References

[1] Pang-Ning Tan et al. *Introduction to data mining*. Pearson Education India, 2006.